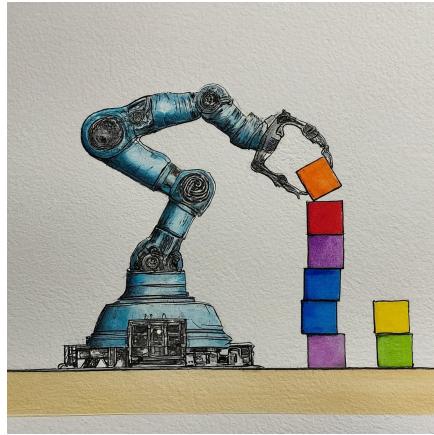


Lab 1 - Introduction to Task Planning

AIR 24-25



Ido Jacobi & Shiran Peeran & Sarah Keren



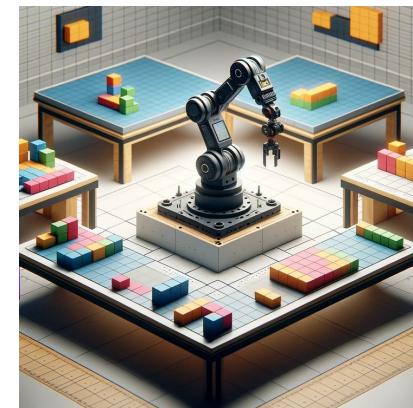
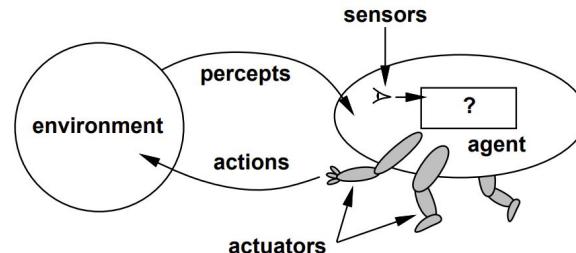
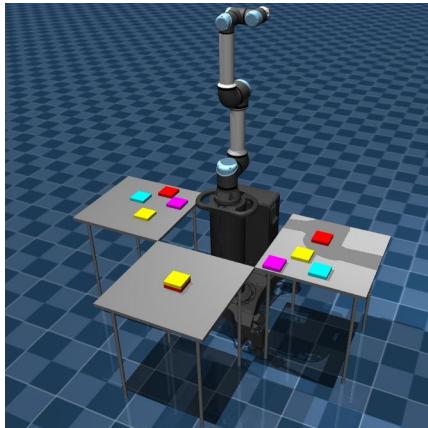
Collaborative AI and Robotics



TECHNION | The Henry and Marilyn Taub
Faculty of Computer Science

Course and Lab Objectives

- The focus of this course: robotic agents that need to perform complex tasks that require
 - high-level reasoning and planning
 - understanding of how to move and operate in the physical environment
- We will start by exploring task planning and explore the limitations of task planning for robotic agents.
- The labs are meant to complement the theoretical material learned in class with in-depth exploration of some of the ideas we learned and with some hands-on experience on getting robots to do what we want them to do



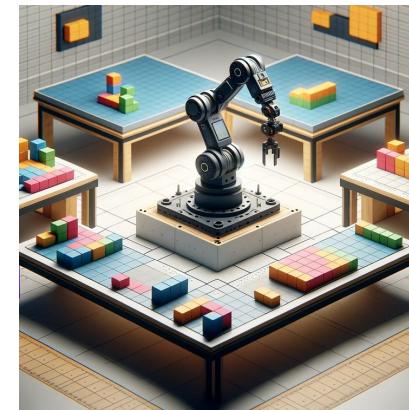
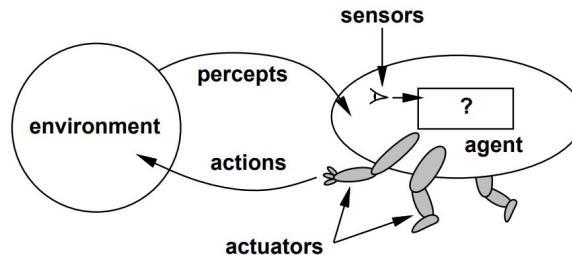
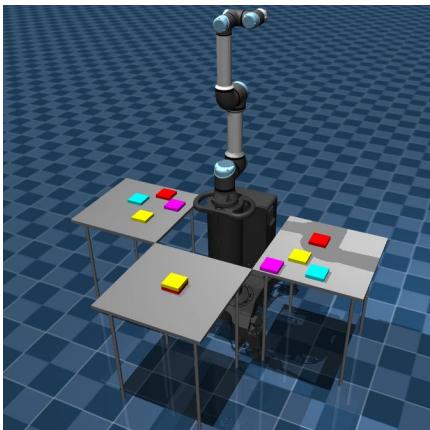
Technical Details

- The tutorials are provided in the form of Jupyter notebooks, suitable for running online via Google Colab.
- Basic knowledge of Python and PDDL is required.
- At the beginning of each notebook, there is a link to run it on Colab.
- If you would like to run the notebooks locally on your machine, you can download them, but some installations may be required (e.g., numpy).
- Each lab will include a presentation (like this one) and a complementary Google notebook with which you can practice the learned ideas.

Our repo: <https://github.com/CLAIR-LAB-TECHNION/AIR>

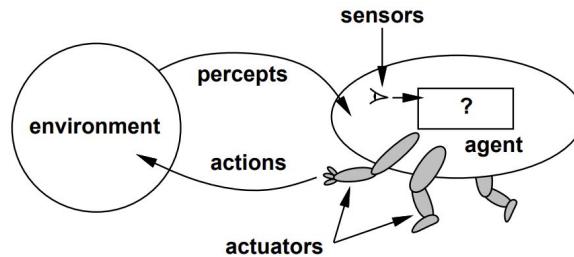
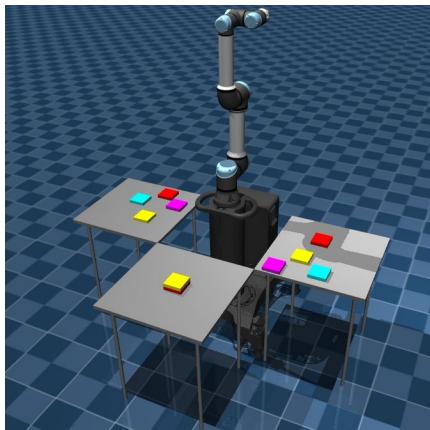
Lab 1 - Objectives

- Getting to know the complexities of task planning for robotic settings
- Understanding PDDL and its limitations of in representing complex settings
- Understanding the need to integrate task and motion planning



Lab 1 - Tasks

1. Getting familiar with PDDL
2. Solving PDDL problems using Best First Search and formulating a heuristic function
3. Understanding the expressive limitations of PDDL



Task Planning with PDDL

Task Planning - Classical Formulation

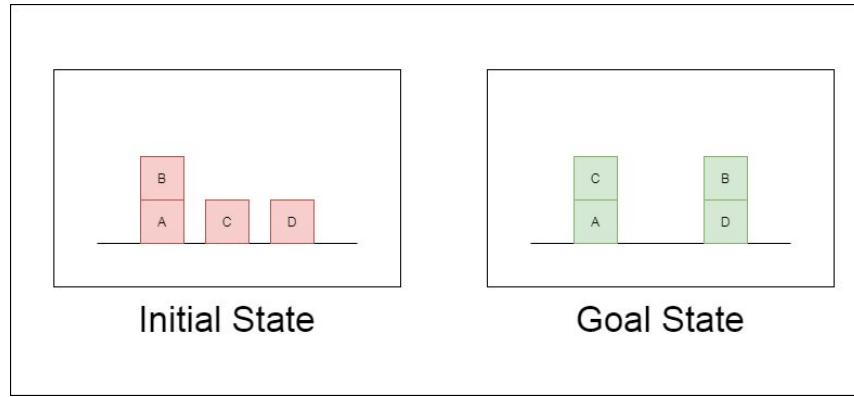
$$\langle S, s_0, S_G, A, f, c \rangle$$

State space induced by feature set V

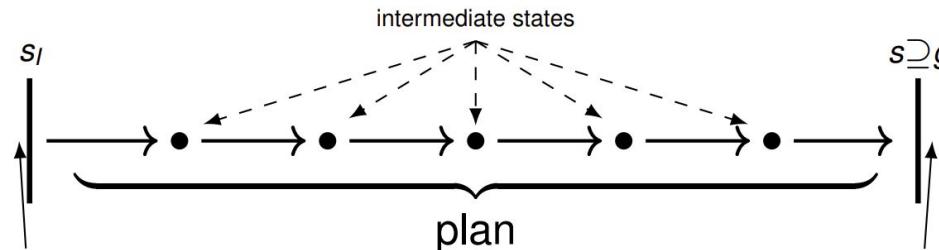
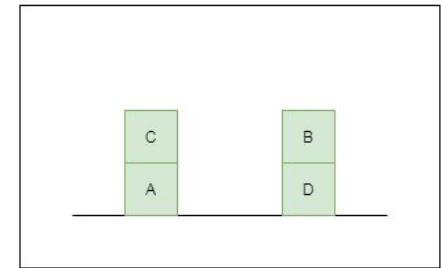
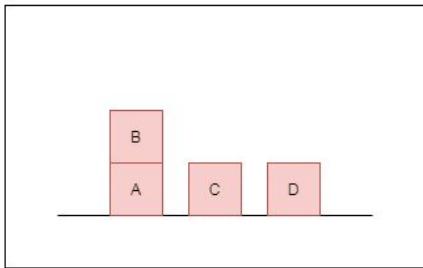
- S is a finite and discrete set of states,
- $s_0 \in S$ is the *known initial state*,
- $S_G \subseteq S$ is the non-empty set of goal states,
- $A(s) \subseteq A$ represents the set of actions in A that are applicable in each state $s \in S$,
- $f(a, s)$ is the *deterministic transition function* where $s' = f(a, s)$ is the state that follows s after doing action $a \in A(s)$, and
- $c(a, s)$ is a *positive cost* for doing action a in the state s .

A solution or *plan* in this model is a sequence of applicable actions a_0, \dots, a_n that generates a state sequence s_0, s_1, \dots, s_{n+1} where s_{n+1} is a goal state. More precisely, the action a_i is applicable

A blocks world example



Planning



description of the
initial world situation

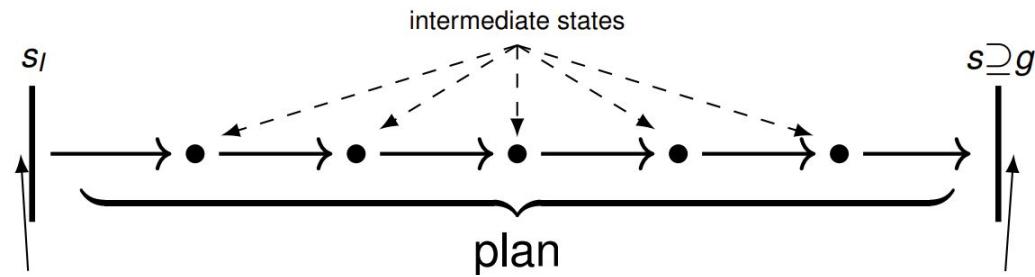
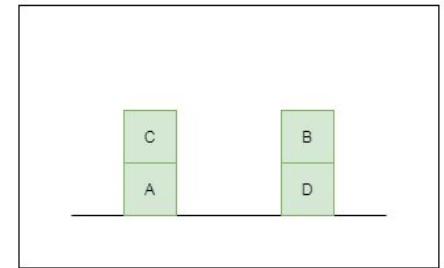
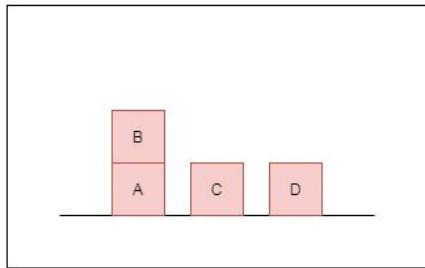
description of desired
world properties

Diagram by
Pascal Bercher

Classical planning: discrete, deterministic, fully observable and single agent.

Problem description induces a state transition system, with nodes as states and actions as transitions.

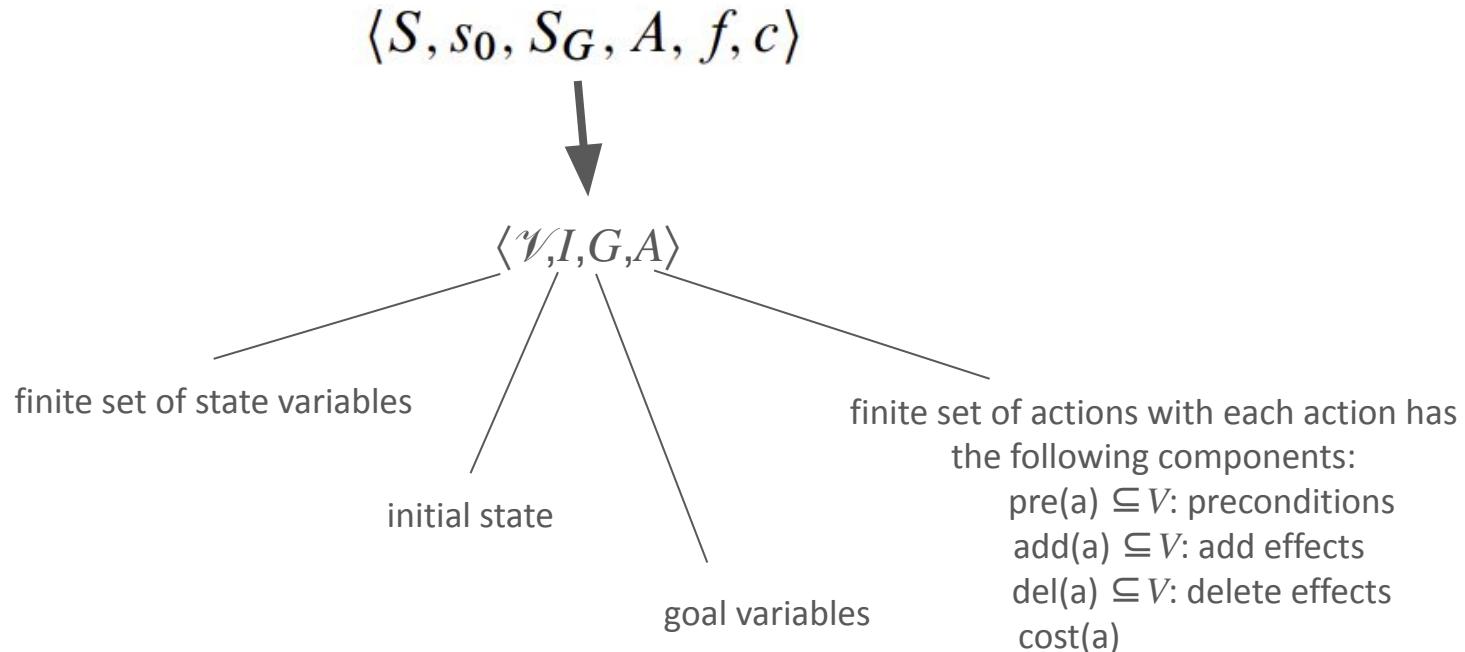
Becomes very complex very quickly: state space is $2^{|V|}$, where V is the number of state features



Pickup B -> Stack B on D -> Pickup C -> Stack C on A



STRIPS and PDDL



STRIPS and PDDL

- STRIPS (Stanford Research Institute Problem Solver) and PDDL (Planning Domain Definition Language) are both used to encoding planning problems.
- PDDL is more expressive:
 - separate domain and problem definitions.
 - domain includes types, predicates, functions, and actions available.
 - The problem description allows describing the objects in the world, the initial state, and the goal.

STRIPS

Operator: MOVE(x, y, z)
Preconditions: ON(x, y), CLEAR(x), CLEAR(z)
Delete list: ON(x, y), CLEAR(z)
Add list: ON(x, z), CLEAR(y)

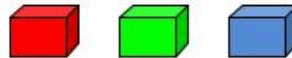
PDDL

```
(:action move
  :parameters (?x ?y ?z - block)
  :precondition (and (on ?x ?y)
                      (clear ?x)
                      (clear ?z))
  :effect (and (not (on ?x ?y))
                (not (clear ?z))
                (on ?x ?z)
                (clear ?y)))
)
```

Formulation of Blocks World

Action/model:

Objects (things):



Predicates (true or false statements about the world, "facts"):

(On A, B)

(Clear A)

Actions (how to change the predicates):

(Clear A)
(Clear B)

Stack

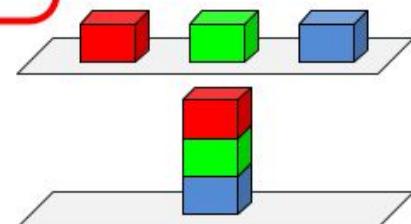
¬(Clear B)
(On A, B)

preconditions

effects

Initial: (Clear Red) (Clear Green) (Clear Blue)

Goal: (On Red, Green) (On Green, Blue)



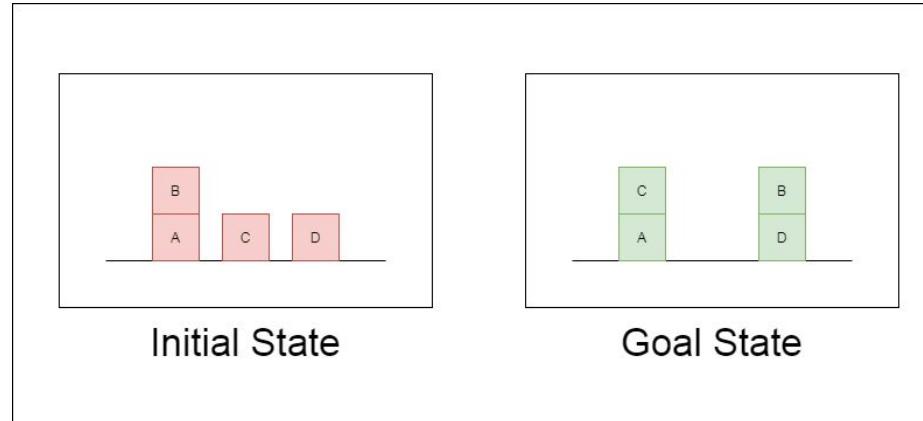
PDDL Formulation of Blocks World

```
(:types block)
(:predicates (on ?b1 ?b2 - block)
             (clear ?b - block)
             (handempty)
             (holding ?b - block))

(:action stack
  :parameters (?b1 ?b2 - block)
  :precondition (and (holding ?b1) (clear ?b2))
  :effect
    (and (not (holding ?b1))
         (not (clear ?b2))
         (clear ?b1)
         (handempty)
         (on ?b1 ?b2)))
```

```
(:init
  (handempty) (ontable D) (clear D) (on B A) ...)

(:goal (and
        (on B D) (ontable D)
        (on C A) (ontable A))))
```



Task 1



Writing a PDDL definition for a
blocks world problem

Task Planning with Best First Search

Task Planning: Solution approaches

Common algorithms/techniques:

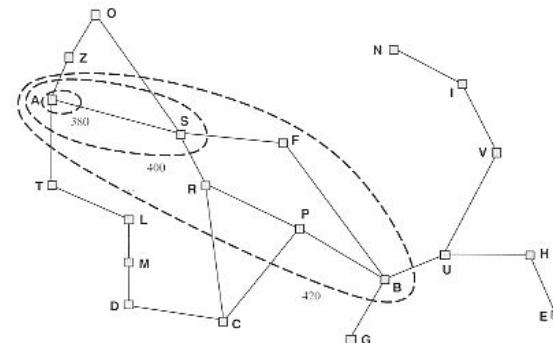
- Search – forward, regression, bidirectional, symbolic
- Compilation – to SAT, to CP
- Partial Order Planning
- Dynamic programming: Bellman Backup – i.e. for problems with uncertainty and rewards, like MDPs and POMDPs.

What are some general search algorithms?

Best First Search

```
function BEST-FIRST-SEARCH(problem, f) returns a solution node or failure
    node  $\leftarrow$  NODE(STATE=problem.INITIAL)
    frontier  $\leftarrow$  a priority queue ordered by f, with node as an element
    reached  $\leftarrow$  a lookup table, with one entry with key problem.INITIAL and value node
    while not IS-EMPTY(frontier) do
        node  $\leftarrow$  POP(frontier)
        if problem.IS-GOAL(node.STATE) then return node
        for each child in EXPAND(problem, node) do
            s  $\leftarrow$  child.STATE
            if s is not in reached or child.PATH-COST < reached[s].PATH-COST then
                reached[s]  $\leftarrow$  child
                add child to frontier
    return failure

function EXPAND(problem, node) yields nodes
    s  $\leftarrow$  node.STATE
    for each action in problem.ACTIONS(s) do
        s'  $\leftarrow$  problem.RESULT(s, action)
        cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s')
        yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```



Code above from Norvig, P. Russel, and S. Artificial Intelligence. "A modern approach." *Prentice Hall Upper Pearl*, Judea. *Heuristics: intelligent search strategies for computer problem solving*. Addison-Wesley Longman Publishing Co., Inc., 1984.

Task Planning - Extensions

Accounting for stochastic action outcome
using a state transition probability function

$$\mathcal{P}_{s,s'}^a = \mathcal{P}[S_{t+1} = s' | S_t = s, A_t = a]$$

Accounting for objectives using a reward function

$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$$

Accounting for reward depreciation using a discount factor

$$\gamma \in [0, 1]$$

Accounting for partial observability using
observations Ω and observation function

$$\mathcal{O}_{s,a}^o = \mathcal{P}[O_{t+1} = o | S_t = s, A_t = a]$$

Accounting for action duration, multiple agents, varying costs, etc.

Sequential decision making: what should the agent do to maximize total rewards?

A*

- A* is a special case of Best First Search that uses a heuristic to guide the search from an initial state to a goal state.
- Formally, the priority according to which A* extracts nodes from the frontier is: $f(n)=g(n)+h(n)$ where:
 - n is the node that is considered,
 - $g(n)$ is the cost of the path from the start node to n
 - $h(n)$ is the heuristic estimation of the cost from n to a goal state.

Heuristic Functions

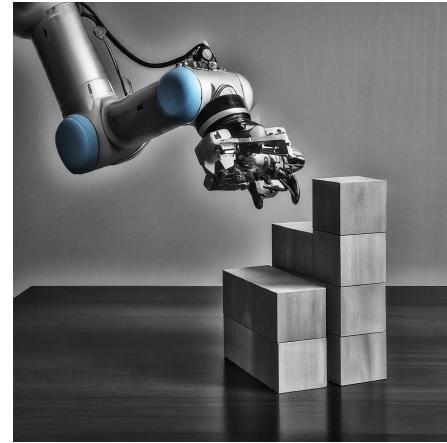
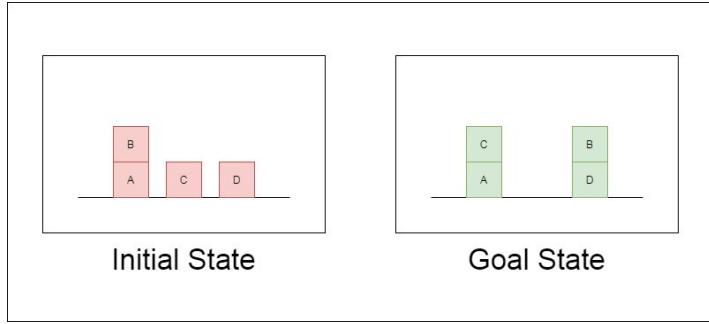
- Heuristic functions are an important element of any heuristic search, and A* in particular.
- Effective heuristics possess certain characteristics that enable the algorithm to efficiently find solutions while minimizing the search effort
- Some key characteristics of heuristic functions:
 - Safeness
 - Goal Awareness
 - Admissibility
 - Efficiency

Task 2



Implementing a heuristic method

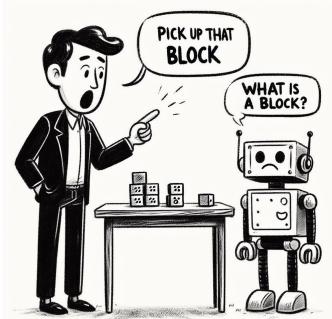
Are we done ?



How does it know what a block is ?

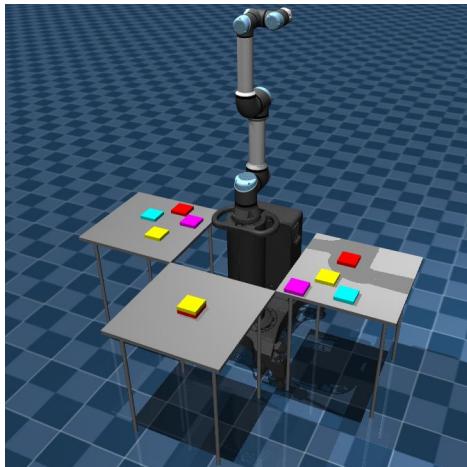
How does it know how to stack, pick up and put down?

How does it know what's the best way to pick up?



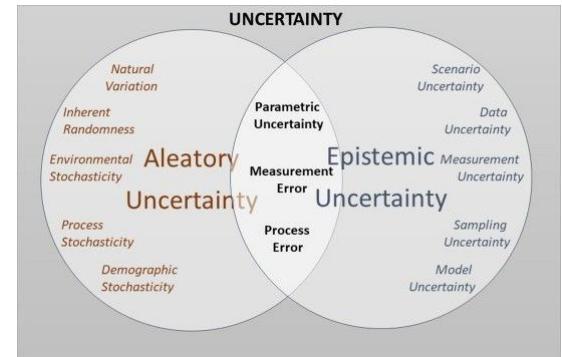
BW in the real(istic) world

What have we failed to capture with our PDDL formulations ?



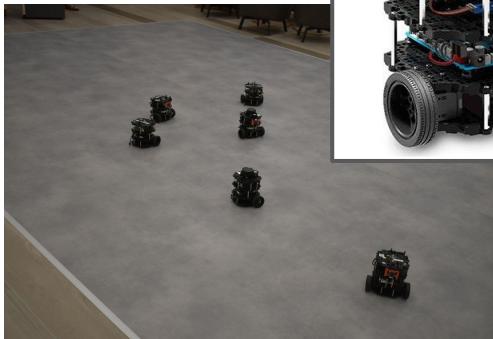
Task Planning vs. the Real World

- **Discrete:** plans are finite sequences of actions - while actual actions are continuous
- **Fully modeled actions** (deterministic in classical planning): while actual actions can have unexpected effects that are hard/impossible to model
- **Full knowledge** of world state or of belief (distribution over possible world states) while knowing the actual state /distribution relies on complex sensing and reasoning



https://www.stat.berkeley.edu/users/aldous/157/Papers/Fox_Ulkumen.pdf

Many Types of Robots



Robot Elements

Sensors



Controllers



Actuators



Robotic Agents



state estimation

Maintain a **belief** $\beta \in \boldsymbol{\beta}$ as a probability distribution over world states S

$$\beta: S \rightarrow [0,1]$$

decision making

Find a **policy**, mapping **belief** β and **objective** into actions (probabilities)

$$\pi: \boldsymbol{\beta} \square A \rightarrow [0,1]$$

motion control

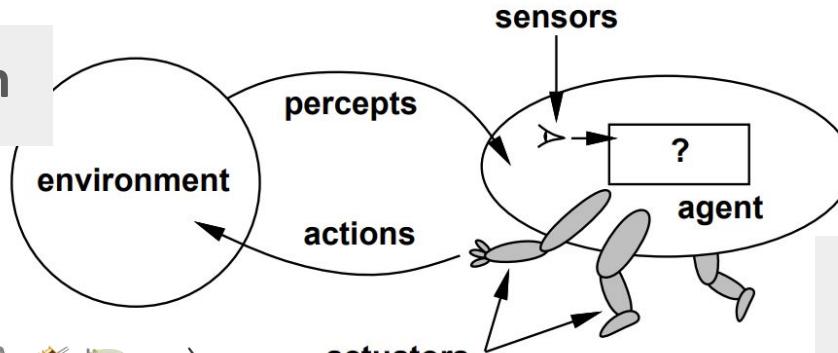
Translate actions into low-level commands (and monitor their execution)

$$u = k_p \cdot e + k_i \cdot \int e dt + k_d \cdot \frac{d}{dt} e$$



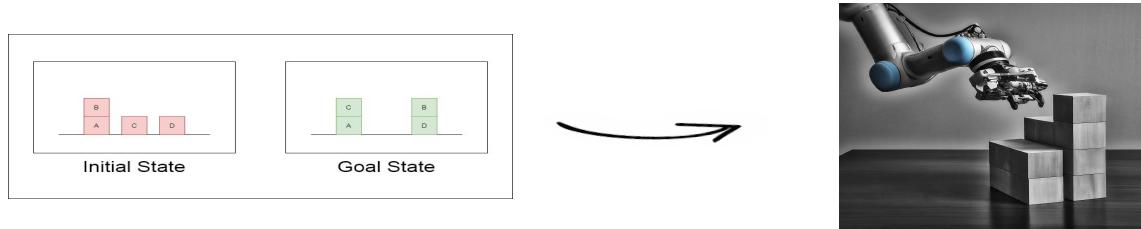
Robotic Agents

state estimation



decision making

motion control



decision making

```
(:action stack
  :parameters (?b1 ?b2 - block)
  :precondition (and (holding ?b1) (clear ?b2))
  :effect
    (and (not (holding ?b1))
         (not (clear ?b2))
         (clear ?b1)
         (handempty)
         (on ?b1 ?b2)))
```

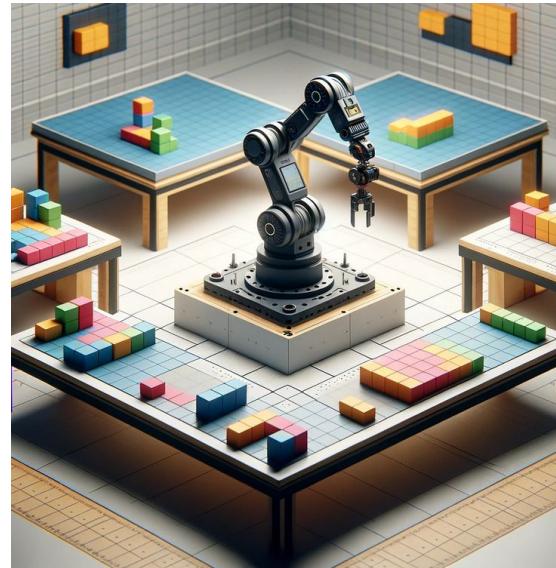
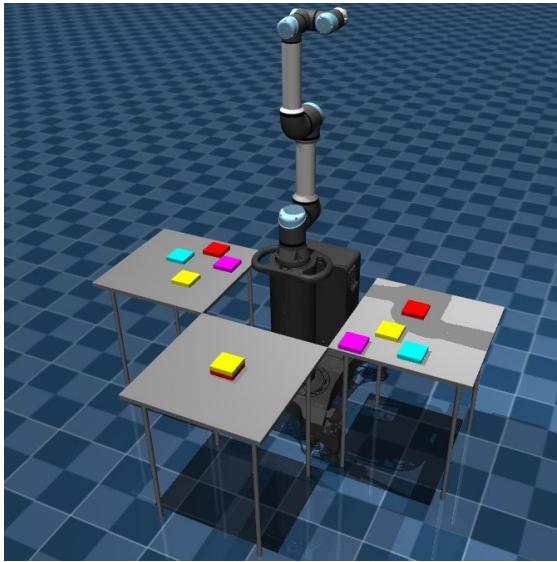
state estimation

motion control

Task planning is not enough

What next ?

Things get messier: N-table Blocks World



Accounting for varying costs, collision avoidance, motion feasibility, multiple (infinite) positions from which an action can be performed, multiple (infinite) ways to reach a position

